

# **CS553 Programming Assignment #1**

## **Benchmarking**

### **Performance Evaluation**

**Harsh C. Parikh  
A20338453**

This document contains all experimental results of all the benchmarks. Each experiment is performed 3 times and average and standard deviation are calculated based on that.

## 1. Introduction

For all the benchmarks, I experimented on Amazon EC2, with 1 core CPU, Intel(R) Xeon(R) CPU E5-2670 v2 @ 2.50GHz. All Programs are done in Python language.

```
ec2-user@ip-172-31-54-135:~  
[ec2-user@ip-172-31-54-135 ~]$ lscpu  
Architecture:          x86_64  
CPU op-mode(s):        32-bit, 64-bit  
Byte Order:            Little Endian  
CPU(s):                1  
On-line CPU(s) list:   0  
Thread(s) per core:    1  
Core(s) per socket:    1  
Socket(s):             1  
NUMA node(s):          1  
Vendor ID:             GenuineIntel  
CPU family:            6  
Model:                 62  
Model name:            Intel(R) Xeon(R) CPU E5-2670 v2 @ 2.50GHz  
Stepping:              4  
CPU MHz:               2494.026  
BogoMIPS:              4988.05  
Hypervisor vendor:     Xen  
Virtualization type:   full  
L1d cache:             32K  
L1i cache:             32K  
L2 cache:              256K  
L3 cache:              25600K  
NUMA node0 CPU(s):    0  
[ec2-user@ip-172-31-54-135 ~]$
```

For CPU benchmark, I measured processor's speed in terms of GFLOPS and GIOPS at varying concurrency level for 1,2 and 4 threads. Moreover, I took the 600 samples for IOPS and FLOPS at the interval of 1 sec till 10 minutes.

For Disk benchmark, I measured the throughputs of Sequential Read, Sequential Write, Random Read and Random Write and latency for all of them and varied the block sizes (1B/1KB/1MB) for the number of threads (1 and 2).

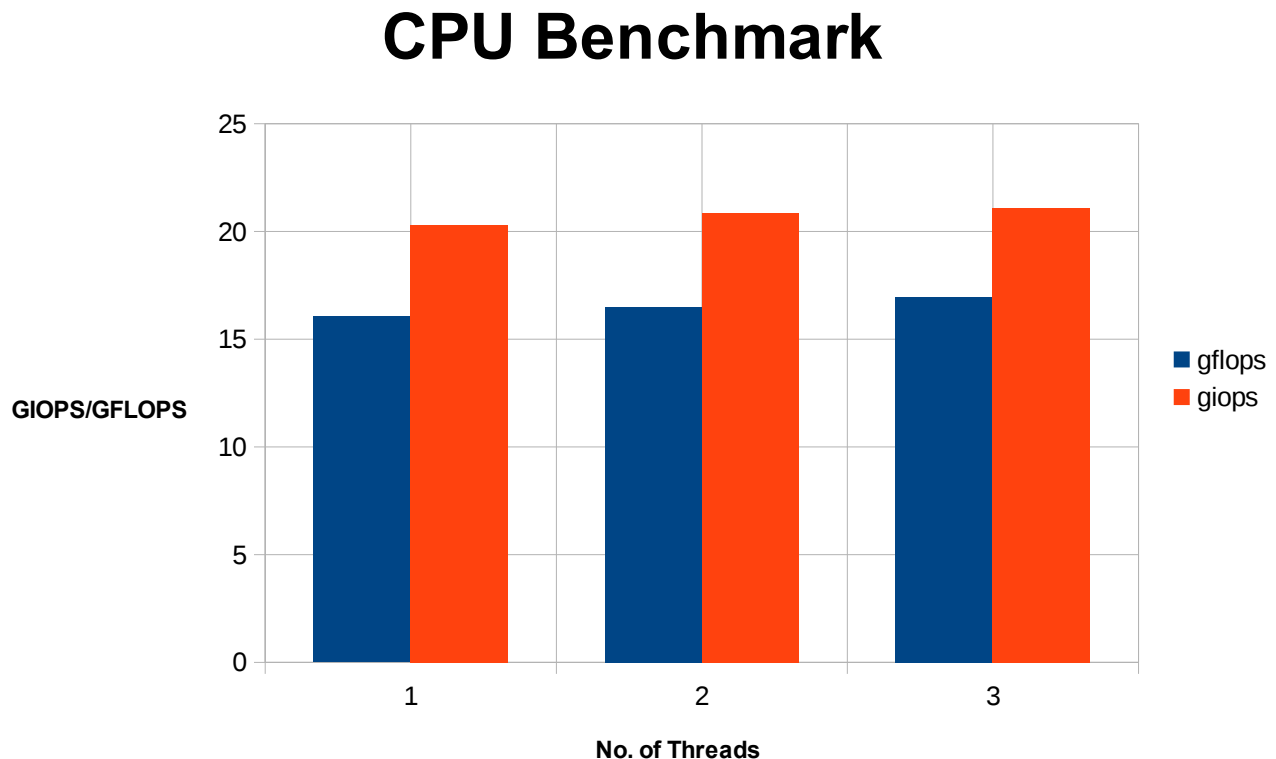
For Network benchmark, I measured throughput and latency for both TCP and UDP protocols and varied the packet size (1B/1KB/64KB) and the number of threads (1 and 2)

## 2. Experiment Results and Analysis

In this section, all experiment results for each benchmark and explanations for the trade offs in the results.

### 2.1 CPU Benchmarking

- For CPU benchmark, I am finding GIOPS and GFLOPS value for different number of threads.
- In graph, I had done multiple experiments for 1,2 and 4 threads and according to that I am getting results for GFLOPS and GIOPS as follows.



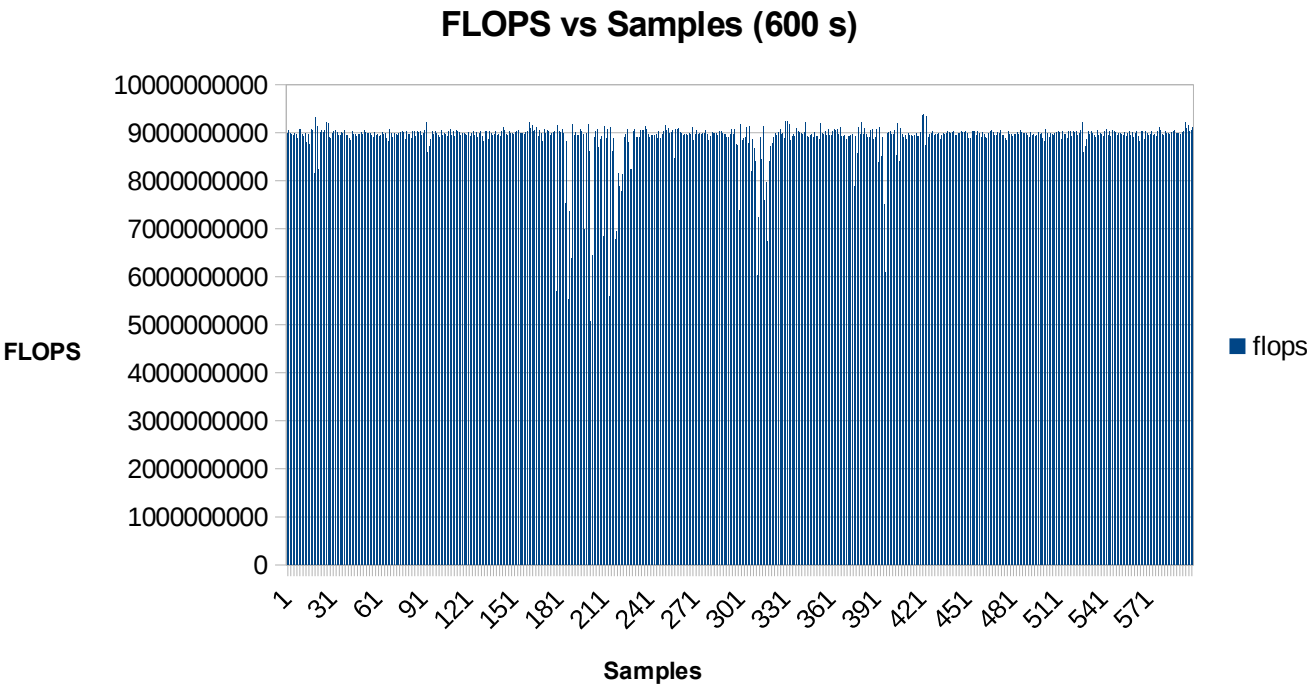
- Theoretical peak performance of the processor is given by:  
$$= (\text{CPU speed in GHz}) \times (\text{number of CPU cores}) \times (\text{CPU instruction per cycle}) \times (\text{number of CPUs per node})$$
$$= 2.50 * 1 * 8 * 2$$
$$= 40 \text{ Gflops}$$

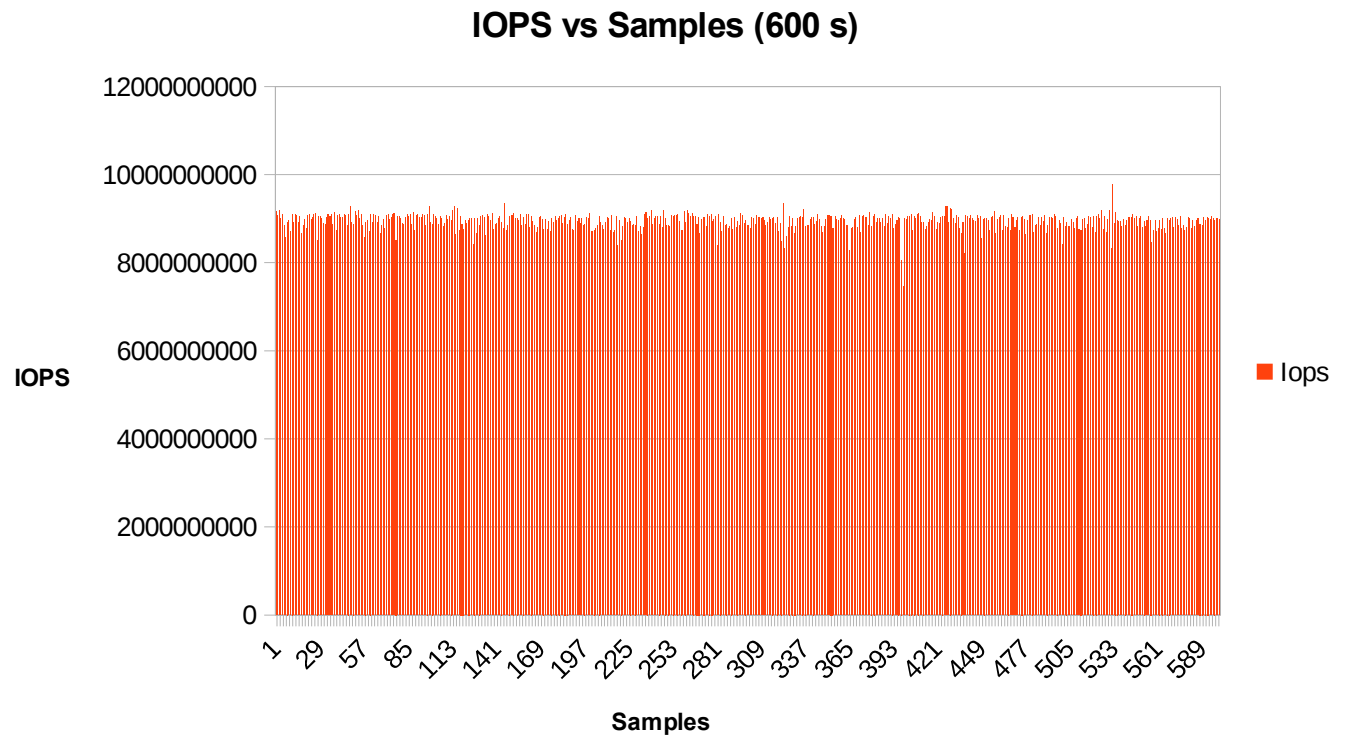
and efficiency would be  $= (20.88/40) * 100 = 52.20 \%$ .

**Average and Standard Deviation for GFLOPS and IOPS are in following tables.**

Number of thread	Average (GIOPS)	Standard Deviation(GIOPS)	Average(GFLOPS)	Standard Deviation(GFLOPS)
1	24.2798	0.3541	20.12567	0.3011
2	24.4782	0.3631	20.8832	0.3166
4	24.9537	0.3777	21.43637	0.3199

Now, I have taken **600 samples of IOPS and FLOPS**. Here are the snap-shots of results. These results performed on **4 threads**.





Here, in the graph, I observe there are some glitches. That's because of operations were performed less in that time interval due to other process might use the core. So sometimes due to that we got these types of glitches as if less numbers of operations performed in that period.

Now, I had run the **Lin-pack benchmark** and result is being shown in below screen..

```
[ec2-user@ip-172-31-59-112 linpack]$ ./runme_xeon64
This is a SAMPLE run script for SMP LINPACK. Change it to reflect
the correct number of CPUs/threads, problem input files, etc..
Thu Feb 11 04:51:24 UTC 2016
Intel(R) Optimized LINPACK Benchmark data

Current date/time: Thu Feb 11 04:51:24 2016

CPU frequency:      2.824 GHz
Number of CPUs: 1
Number of cores: 1
Number of threads: 1

Parameters are set to:

Number of tests: 15
Number of equations to solve (problem size) : 1000  2000  5000  10000 15000 18000
0 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array      : 1000  2000  5008  10000 15000 18000
8 20016 22008 25000 26000 27000 30000 35000 40000 45000
Number of trials to run        : 4      2      2      2      2      2
2      2      2      2      1      1      1      1      1      1
Data alignment value (in Kbytes) : 4      4      4      4      4      4
4      4      4      4      4      1      1      1      1      1

Maximum memory requested that can be used=800204096, at the size=10000

===== Timing linear equation system solver =====

Size   LDA   Align. Time(s)   GFlops   Residual   Residual(norm) Check
1000   1000   4       0.040   16.6632   9.900691e-13 3.376390e-02 pass
1000   1000   4       0.038   17.6646   9.900691e-13 3.376390e-02 pass
1000   1000   4       0.039   17.1456   9.900691e-13 3.376390e-02 pass
1000   1000   4       0.038   17.6622   9.900691e-13 3.376390e-02 pass
2000   2000   4       0.286   18.6465   4.053480e-12 3.526031e-02 pass
2000   2000   4       0.286   18.6809   4.053480e-12 3.526031e-02 pass
5000   5008   4       4.263   19.5591   2.336047e-11 3.257429e-02 pass
5000   5008   4       4.265   19.5509   2.336047e-11 3.257429e-02 pass
10000  10000  4       33.075   20.1620   1.124127e-10 3.963786e-02 pass
10000  10000  4       32.466   20.5406   1.124127e-10 3.963786e-02 pass

Performance Summary (GFlops)

Size   LDA   Align. Average Maximal
1000   1000   4       17.2839 17.6646
2000   2000   4       18.6637 18.6809
5000   5008   4       19.5550 19.5591
10000  10000  4       20.3513 20.5406

Residual checks PASSED

End of tests

Done: Thu Feb 11 04:52:46 UTC 2016
[ec2-user@ip-172-31-59-112 linpack]$
```

From practical performance, linpack gives me around **20.35 Gflops**.

**Efficiency = (20.53/40)\*100 = 51.32 %**

## 2.2 Disk Benchmark

```
iozone: Performance Test of File I/O
Version $Revision: 3.420 $
Compiled for 64 bit mode.
Build: linux-AMD64

Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collins
              Al Slater, Scott Rhine, Mike Wisner, Ken Goss
              Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
              Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
              Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
              Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
              Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
              Vangel Bojaxhi, Ben England, Vikentsi Lapa.

Run began: Fri Feb 12 15:10:01 2016

File size set to 1024 KB
Command line used: iozone -i -a -i 0 -i 1 -i 2 -s 1024
Output is in Kbytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 Kbytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

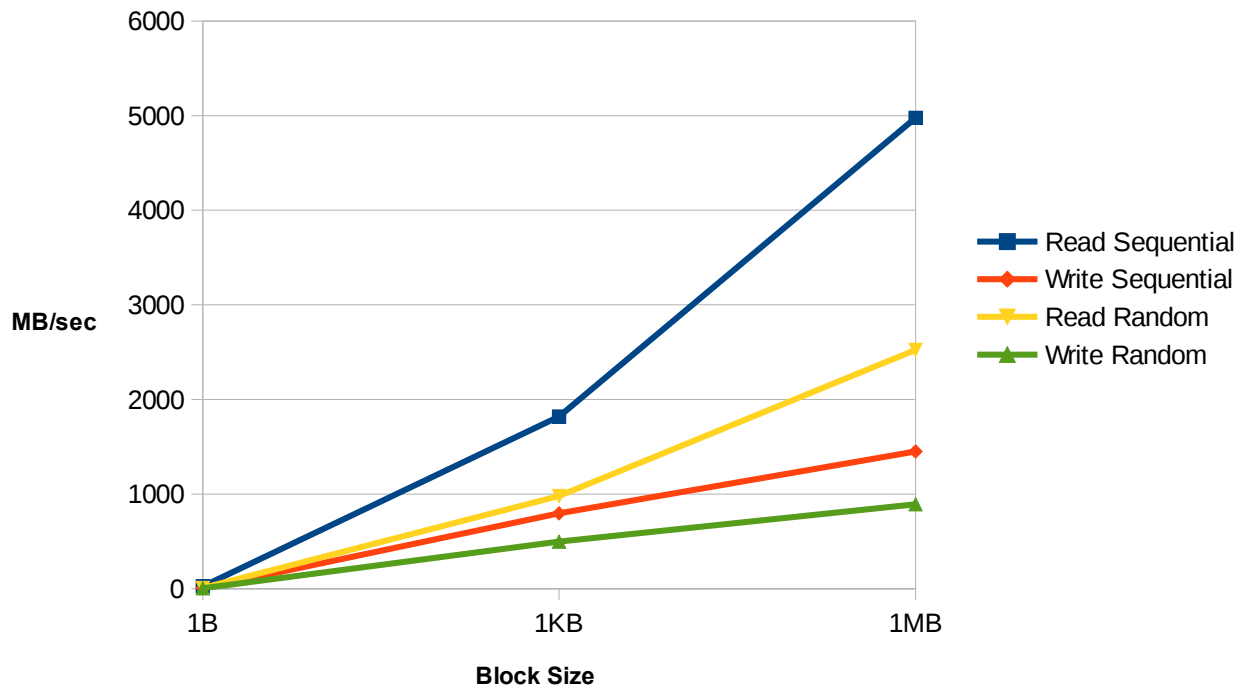
      random  random  bkwd  record  stride
KB  reflen  write rewrite  read  reread  read  write  read  rewrite  read  fwrite frewrite  fread  freread
1024    4 1261072 2326879 4995276 6317933 6402699 2666347

iozone test complete.
```

From practical Throughput for reading 1MB or 1024 KB, I'm getting around = 4.977 GB/sec = 4977615 B/sec and as per theoretical Throughput, it would be around 6 GB/sec (as per wikipedia, our EC2 instance version is 3.0) . So efficiency would be around **82.83%**. Moreover on IOzone, throughput is around 6.40 GB/sec.

- And latency is = 0.747 ms .
- Here I have calculated the throughput of single thread and multi threaded and block size for sequential and random read/write. And their results are shown.
- Here, I'm getting less throughput for 2 threads compare to 1 thread because when in 2 threads, after completion of one thread, other can access the file. So 2 threads can't access a file at the same time. (Read it somewhere from Amazon Notes).

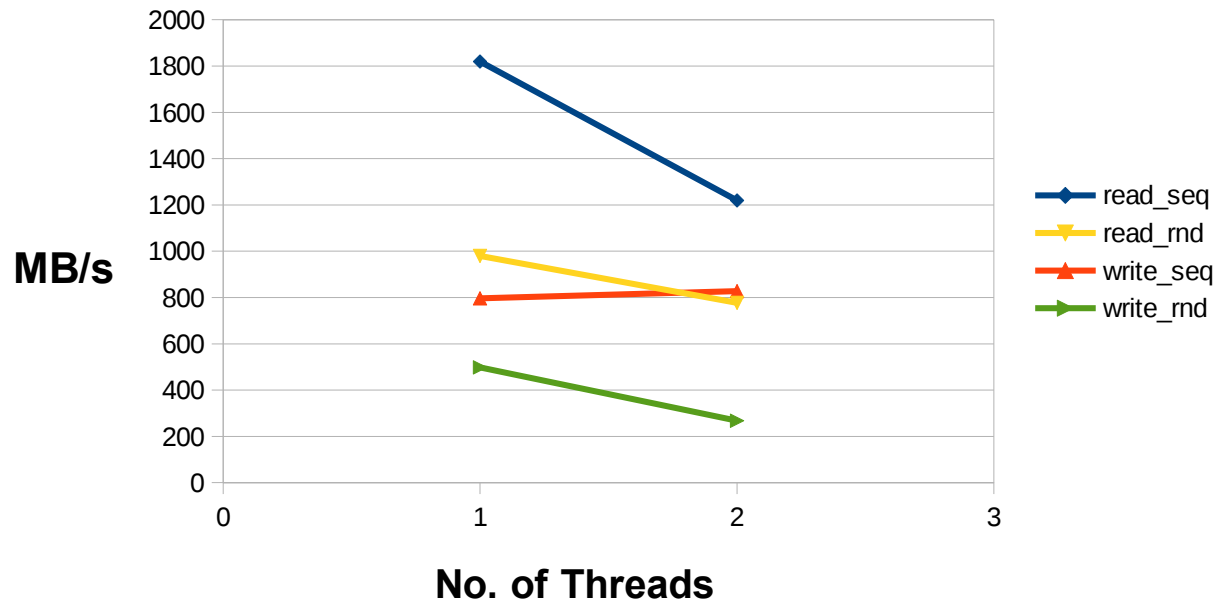
## Throughput (1 Thread)



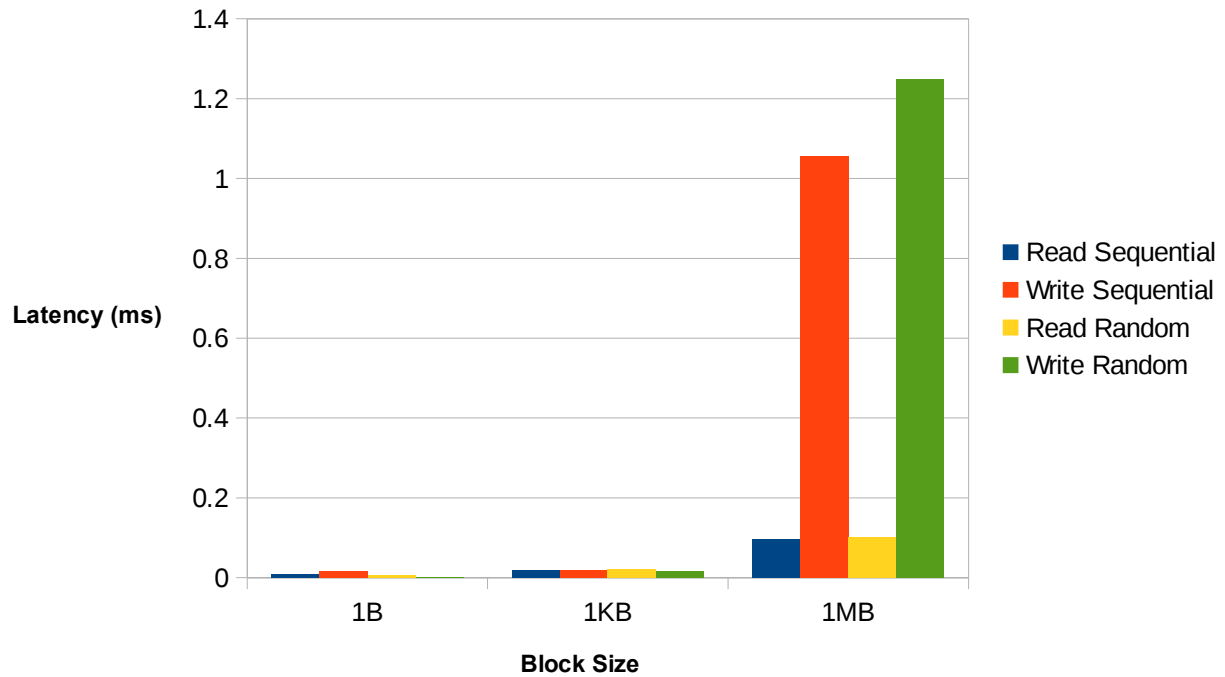
- The graph shows the trade offs between sequential and random read/write throughput according to block size. It's going upwards as we increase the block sizes. And if we increase the number of threads, then throughput is sometimes decreasing because if one can't write in a file simultaneously. So that
- Sequential throughput is faster than random throughput for both read and write operations.
- Read throughput performs better in both sequential and random cases.
- Another graph shows the trade offs between the number of threads and throughput in sequential read, sequential write, random read, and random write with respect to latency in milliseconds.
- Sequential operation works faster than random throughput for both read and write.
- In other graphs, I have used 1 KB block to measure the trade-off between latency and thread. So, one can understand from the graph that if a multi-threaded operation takes less time for both of them.



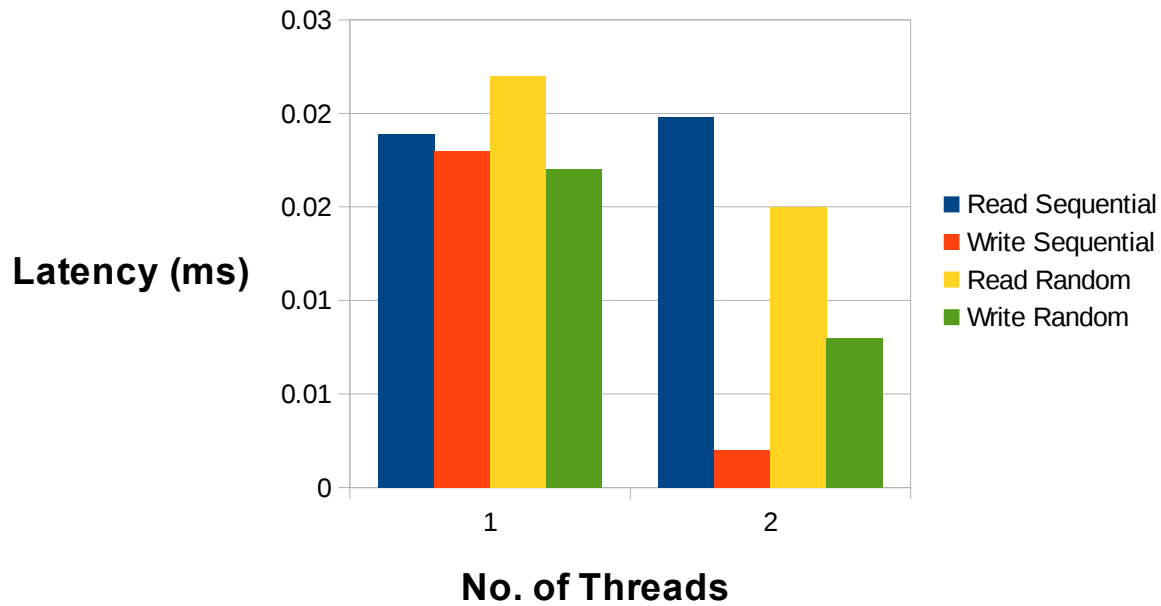
# Throughput(1 KB)



## Latency vs Block size



## Latency vs Thread



### Throughput result of Sequential Read/Write Operations (results in MB/sec):

Threads	Read			Write		
	1b	1k	1m	1b	1k	1m
1	24.138	1819.839	4977.615	15.442	796.307	1451.383
2	22.339	1219.013	3988.177	18.996	827.156	1290.241

### Random Read/Write Operations (results in MB/sec):

Threads	Read			Write		
	1b	1k	1m	1b	1k	1m
1	18.45	979.985	2523.345	14.998	497.897	892.291
2	20.562	776.589	2069.252	15.945	267.586	579.468

**Average and Standard Deviation for 1 Byte are in following tables(results in MB/sec & latency in ms).**

**For Sequential Read (1 B):**

<b>Number of Thread</b>	<b>Average(Throughput-Sequential Read)</b>	<b>Standard Deviation(Throughput-Sequential Read)</b>	<b>Average(Latency-Sequential Read)</b>	<b>Standard Deviation((Latency-Sequential Read)</b>
1	24.138	4.22	0.0043	0.0036
2	22.339	3.88	0.0042	0.0027

**For Random read (1 B):**

<b>Number of Thread</b>	<b>Average(Throughput-Random Read)</b>	<b>Standard Deviation(Throughput-Random Read)</b>	<b>Average(Latency-Random Read)</b>	<b>Standard Deviation((Latency-Random Read)</b>
1	18.45	6.99	0.0163	0.006
2	20.562	5.29	0.0071	0.004

**For Sequential write (1 B):**

<b>Number of Thread</b>	<b>Average(Throughput-Sequential Write)</b>	<b>Standard Deviation(Throughput-Sequential Write)</b>	<b>Average(Latency-Sequential write)</b>	<b>Standard Deviation((Latency-Sequential Write)</b>
1	15.442	2.88	0.081	0.007
2	18.996	4.22	0.096	0.006

**For Random write (1 B):**

<b>Number of Thread</b>	<b>Average(Throughput-Random write)</b>	<b>Standard Deviation(Throughput-Random write)</b>	<b>Average(Latency-Random write)</b>	<b>Standard Deviation((Latency-Random write)</b>
1	14.998	7.12	0.0101	0.006
2	15.545	3.65	0.0123	0.008

```
iozone: Performance Test of File I/O
Version $Revision: 3.420 $
Compiled for 64 bit mode.
Build: linux-AMD64

Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collins
Al Slater, Scott Rhine, Mike Wisner, Ken Goss
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
Vangel Bojaxhi, Ben England, Vikentsi Lapa.

Run began: Fri Feb 12 15:10:01 2016

File size set to 1024 KB
Command line used: iozone -i -a -i 0 -i 1 -i 2 -s 1024
Output is in Kbytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 Kbytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

      random random  bkwd  record  stride
KB reflen  write rewrite  read  reread  read  write  read  rewrite  read  fwrite frewrite  fread  freread
1024      4 1261072 2326879 4995276 6317933 6402699 2666347

iozone test complete.
```

## 2.3 Network Benchmark

- The following graph shows the trade offs between Latency and throughput according to block size. It's going upwards as we increase the block sizes in both TCP and UDP .
- In Network, I observe that TCP throughput is lesser than UDP because UDP is a connectionless protocol and so in UDP, there is no guarantee of receiving bytes.
- On the other side, TCP sends an acknowledgement from server. So for that reason, in TCP, we are getting higher latency compare to UDP.
- The graph shows the trade offs between Latency and throughput according to block size. It's going upwards as we increase the block sizes in both TCP and UDP .
- In Network, I observe that TCP throughput is lesser than UDP because UDP is a connectionless protocol and so in UDP, there is no guarantee of receiving packages.
- On the other side, TCP sends an acknowledgement from server. So for that reason, in TCP, we are getting higher latency compare to UDP.

**These are two snap-shots of UDP and TCP Iperf performance.**

**UDP Graph: -**

```
r@ip-172-31-54-135:~
[ec2-user@ip-172-31-54-135 ~]$ sudo iperf3 -c ec2-54-86-192-251.compute-1.amazonaws.com -p 80 -u -b 100m
Connecting to host ec2-54-86-192-251.compute-1.amazonaws.com, port 80
[ 4] local 172.31.54.135 port 48724 connected to 172.31.54.135 port 80
[ ID] Interval      Transfer    Bandwidth  Total Datagrams
[ 4] 0.00-1.00 sec  10.8 MBytes 90.2 Mbits/sec 1377
[ 4] 1.00-2.00 sec  11.9 MBytes 99.9 Mbits/sec 1525
[ 4] 2.00-3.00 sec  11.9 MBytes 100 Mbits/sec 1526
[ 4] 3.00-4.00 sec  11.9 MBytes 100 Mbits/sec 1526
[ 4] 4.00-5.00 sec  11.9 MBytes 100 Mbits/sec 1526
[ 4] 5.00-6.00 sec  11.9 MBytes 100 Mbits/sec 1526
[ 4] 6.00-7.00 sec  11.9 MBytes 100 Mbits/sec 1526
[ 4] 7.00-8.00 sec  11.9 MBytes 100 Mbits/sec 1526
[ 4] 8.00-9.00 sec  11.9 MBytes 100 Mbits/sec 1526
[ 4] 9.00-10.00 sec 11.9 MBytes 99.9 Mbits/sec 1525
-----
[ ID] Interval      Transfer    Bandwidth  Jitter    Lost/Total Datagrams
[ 4] 0.00-10.00 sec 118 MBytes 99.0 Mbits/sec 0.003 ms  0/15109 (0%)
[ 4] Sent 15109 datagrams

iperf Done.
[ec2-user@ip-172-31-54-135 ~]$
```

Here, in this Iperf, they have given Bandwidth and packets those are being transferred from transmitter to receiver side. In UDP, we can't send more than 65536 bytes package at the same time.

Moreover, in UDP, we need to send packets one by one and in TCP, we can send number of bytes together. So in screen-shot, we get the idea about that.

Efficiency comparison between Iperf and My result.

$$\text{Efficiency} = (759/979) * 100 = 77.52\%$$

## TCP Graph:

```
ec2-user@ip-10-0-1-216 ~]$ sudo iperf3 -c ec2-54-164-252-91.compute-1.amazonaws.com -t 1 -t 60 -V -p 80
iperf 3.0.11
Linux ip-10-0-1-216 4.1.10-17.31.anzn1.x86_64 #1 SMP Sat Oct 24 01:31:37 UTC 2015 x86_64 x86_64 x86_64 GNU/Linux
Time: Fri, 12 Feb 2016 05:52:39 GMT
Connecting to host ec2-54-164-252-91.compute-1.amazonaws.com, port 80
Cookie: ip-10-0-1-216.1455256350.414578.22ff
TCP MSS: 8949 (default)
[ 4] local 10.0.1.216 port 39189 connected to 10.0.0.59 port 80
Starting Test: protocol: TCP, 1 streams, 131072 byte blocks, omitting 0 seconds, 60 second test
[ ID] Interval      Transfer      Bandwidth      Retr  Cwnd
[ 4] 0.00-1.00 sec    105 MBytes    879 Mbits/sec    13    821 KBytes
[ 4] 1.00-2.00 sec    96.2 MBytes    807 Mbits/sec    12    743 KBytes
[ 4] 2.00-3.00 sec    110 MBytes    923 Mbits/sec     0    839 KBytes
[ 4] 3.00-4.00 sec    105 MBytes    881 Mbits/sec     0    874 KBytes
[ 4] 4.00-5.00 sec    81.2 MBytes    682 Mbits/sec     0    874 KBytes
[ 4] 5.00-6.00 sec    82.5 MBytes    692 Mbits/sec     0    874 KBytes
[ 4] 6.00-7.00 sec    82.5 MBytes    692 Mbits/sec     0   1005 KBytes
[ 4] 7.00-8.00 sec    82.5 MBytes    692 Mbits/sec     0   1005 KBytes
[ 4] 8.00-9.00 sec    83.8 MBytes    703 Mbits/sec     0   1005 KBytes
[ 4] 9.00-10.00 sec   82.5 MBytes    692 Mbits/sec     0   1005 KBytes
[ 4] 10.00-11.00 sec   83.8 MBytes    703 Mbits/sec     0    1.55 MBytes
[ 4] 11.00-12.00 sec   83.8 MBytes    703 Mbits/sec     0    1.55 MBytes
[ 4] 12.00-13.00 sec   83.8 MBytes    703 Mbits/sec     0    1.55 MBytes
[ 4] 13.00-14.00 sec   97.5 MBytes    810 Mbits/sec     8    813 KBytes
[ 4] 14.00-15.00 sec   111 MBytes    933 Mbits/sec     0    979 KBytes
[ 4] 15.00-16.00 sec   110 MBytes    923 Mbits/sec     8    935 KBytes
[ 4] 16.00-17.00 sec   109 MBytes    912 Mbits/sec     5    769 KBytes
[ 4] 17.00-18.00 sec   106 MBytes    891 Mbits/sec     7    821 KBytes
[ 4] 18.00-19.00 sec   110 MBytes    923 Mbits/sec    12    769 KBytes
[ 4] 19.00-20.00 sec   110 MBytes    923 Mbits/sec     4    752 KBytes
[ 4] 20.00-21.00 sec   108 MBytes    902 Mbits/sec     6    900 KBytes
[ 4] 21.00-22.00 sec   93.8 MBytes    786 Mbits/sec     2    961 KBytes
[ 4] 22.00-23.00 sec   65.0 MBytes    545 Mbits/sec     0    1.01 MBytes
[ 4] 23.00-24.00 sec   105 MBytes    881 Mbits/sec     3    918 KBytes
[ 4] 24.00-25.00 sec   109 MBytes    912 Mbits/sec    11    900 KBytes
[ 4] 25.00-26.00 sec   102 MBytes    860 Mbits/sec     8    708 KBytes
[ 4] 26.00-27.00 sec   86.2 MBytes    724 Mbits/sec     2    900 KBytes
[ 4] 27.00-28.00 sec   105 MBytes    881 Mbits/sec     3    821 KBytes
[ 4] 28.00-29.00 sec   110 MBytes    923 Mbits/sec    17    856 KBytes
[ 4] 29.00-30.00 sec   108 MBytes    902 Mbits/sec     5    935 KBytes
[ 4] 30.00-31.00 sec   88.8 MBytes    744 Mbits/sec     0    961 KBytes
[ 4] 31.00-32.00 sec   88.8 MBytes    744 Mbits/sec     3    699 KBytes
[ 4] 32.00-33.00 sec   92.5 MBytes    776 Mbits/sec     0    848 KBytes
[ 4] 33.00-34.00 sec   98.8 MBytes    829 Mbits/sec     7    918 KBytes
[ 4] 34.00-35.00 sec   87.5 MBytes    734 Mbits/sec     5    848 KBytes
[ 4] 35.00-36.00 sec   106 MBytes    891 Mbits/sec     5    830 KBytes
[ 4] 36.00-37.00 sec   105 MBytes    881 Mbits/sec    10    725 KBytes
[ 4] 37.00-38.00 sec   105 MBytes    881 Mbits/sec     4    673 KBytes
[ 4] 38.00-39.00 sec   109 MBytes    912 Mbits/sec     4    848 KBytes
[ 4] 39.00-40.00 sec   105 MBytes    881 Mbits/sec     4    970 KBytes
[ 4] 40.00-41.00 sec   109 MBytes    912 Mbits/sec     0    1.05 MBytes
[ 4] 41.00-42.00 sec   102 MBytes    860 Mbits/sec    19    944 KBytes
[ 4] 42.00-43.00 sec   106 MBytes    891 Mbits/sec     2    961 KBytes
[ 4] 43.00-44.00 sec   102 MBytes    860 Mbits/sec     5    848 KBytes
[ 4] 44.00-45.00 sec   101 MBytes    849 Mbits/sec     3    961 KBytes
[ 4] 45.00-46.00 sec   102 MBytes    860 Mbits/sec     8    778 KBytes
[ 4] 46.00-47.00 sec   106 MBytes    891 Mbits/sec     6    988 KBytes
[ 4] 47.00-48.00 sec   108 MBytes    902 Mbits/sec     0    1.03 MBytes
[ 4] 48.00-49.00 sec   109 MBytes    912 Mbits/sec     0    1.06 MBytes
[ 4] 49.00-50.00 sec   108 MBytes    902 Mbits/sec    11    801 KBytes
[ 4] 50.00-51.00 sec   108 MBytes    902 Mbits/sec    15    918 KBytes
[ 4] 51.00-52.00 sec   96.2 MBytes    807 Mbits/sec     3    821 KBytes
[ 4] 52.00-53.00 sec   90.0 MBytes    755 Mbits/sec     3    690 KBytes
[ 4] 53.00-54.00 sec   110 MBytes    923 Mbits/sec    15    568 KBytes
[ 4] 54.00-55.00 sec   109 MBytes    912 Mbits/sec     5    734 KBytes
[ 4] 55.00-56.00 sec   101 MBytes    849 Mbits/sec     2    926 KBytes
[ 4] 56.00-57.00 sec   92.5 MBytes    776 Mbits/sec     1    865 KBytes
[ 4] 57.00-58.00 sec   98.8 MBytes    755 Mbits/sec    22    664 KBytes
[ 4] 58.00-59.00 sec   102 MBytes    860 Mbits/sec    18    673 KBytes
[ 4] 59.00-60.00 sec   98.8 MBytes    828 Mbits/sec     5    638 KBytes
-----
Test Complete. Summary Results:
[ ID] Interval      Transfer      Bandwidth      Retr
[ 4] 0.00-60.00 sec  5.81 GBytes    831 Mbits/sec    311      sender
[ 4] 0.00-60.00 sec  5.80 GBytes    831 Mbits/sec      receiver
CPU Utilization: local/sender 2.3% (0.0%u/2.3%k), remote/receiver 6.4% (0.8%u/5.6%k)

iperf Done.
ec2-user@ip-10-0-1-216 ~]$
```

### Throughput (1 Thread) in MB/sec:

Number of blocks	Average(Throughput-TCP)	Standard Deviation(Throughput-TCP)	Average(Throughput-UDP)	Standard Deviation(Throughput-UDP)
1 B	1.023	0.324	2.085	1.1445
1 KB	328.614	27.76	452.225	29.897
64 KB	759.314	85.14	1310.124	110.255

### Latency (1 Thread) in ms:

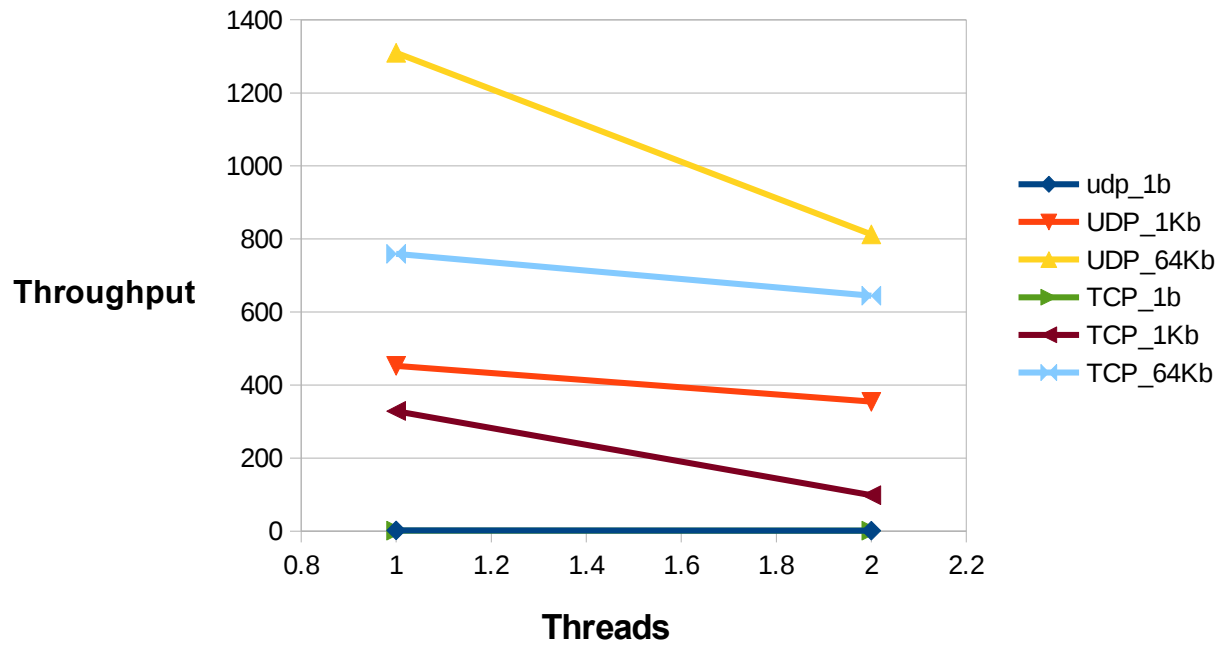
Number of blocks	Average(Latency-TCP)	Standard Deviation(Latency-TCP)	Average(Latency-UDP)	Standard Deviation(Latency-UDP)
1 B	0.147	0.005	0.077	0.0025
1 KB	0.506	0.002	0.498	0.05
64 KB	12.372	1.992	8.0079	1.56

### Throughput result of TCP and UDP (results in MB):

UDP				
Thread	1b	1Kb	64Kb	
1	2.023	452.25	1310.14	
2	1.085	354.233	812.5	

TCP				
Thread	1b	1Kb	64Kb	
1	1.023	328.614	759.314	
2	1.069	98.01	644.259	

## Throughput vs Threads



## Throughput vs Block

